
The Multivac: Blind Peer Matrix Evaluation of Frontier Language Models

Yash Darji
Independent Researcher
ORCID: [0009-0009-6895-842X](https://orcid.org/0009-0009-6895-842X)
yashdarji2378@gmail.com

Abstract

Single-judge evaluation of large language models introduces systematic bias (Zheng et al., 2023). We propose the **Blind Peer Matrix**, a multi-judge evaluation methodology in which N frontier models both generate responses and evaluate each other’s outputs in a fully blinded $N \times N$ matrix, with self-judgments excluded. We apply this method across 286 evaluations spanning 198 unique questions in 9 category pools (code generation, reasoning, analysis, communication, meta-alignment, edge cases, and three focused model-family batches), producing 22,252 valid scored judgments (23,356 parsed) from 55 model configurations (50 distinct by name) across 17 vendor families. Three findings stand out. First, no category has a statistically separated winner: in all nine pools the top model beats second place in only 50–58% of matched comparisons, with the difference interval spanning 50%. Pooling categories supplies the power a single 30-question pool lacks—among broad-participation generalists GPT-5.4 holds rank 1 in 99% of question-clustered resamples and is the only rank-1 coverage member, though Grok 4.20 is not excludable under family-wise correction and GPT-5.4’s edge over it is a 54% win rate. The identity of the “overall” champion is therefore an artifact of the aggregation: naive mean crowns leniency-favored models and a specialist-polluted global ranking a near coin-flip, and only leniency-adjusted and exposure-restricted pooling yields GPT-5.4. Second, and centrally, we show that the standard estimator for same-vendor judge bias is confounded by respondent quality and judge leniency. Decomposing all eight testable families algebraically and re-estimating with a within-response fixed-effects model with standard errors clustered two-way by judge and question, only two families show same-vendor favoritism that survives the controls—Anthropic (+0.41, robust: worst-case leave-one-judge-out $p < 10^{-4}$) and MiniMax (+0.40, significant under the primary specification at $p = 0.001$ but judge-dependent—its two-way significance does not survive leave-one-judge-out, worst-case $p \approx 0.02$)—while Qwen (+0.56) does not survive the two-way clustering. No family shows significant negative bias: the naive estimator’s strong negatives for Mistral (−1.02) and Google (−0.59), and its positive for xAI (+0.75), are artifacts of judge leniency and respondent quality that do not survive controls. The same pattern—a routine estimator confounded and exposed by counterexample—recurs in aggregate scoring itself: silently discarding non-responses, standard practice in leaderboard construction, inflates the mean of models that frequently refuse by up to four points on a ten-point scale, systematically favoring them. Third, judge disagreement is category-dependent, with code evaluation producing just over double the disagreement of meta-alignment tasks ($\sigma = 1.27$ vs. 0.63). We release the complete evaluation dataset (27,540 judgments including self-exclusions), an open-source evaluation framework, all question prompts, and a one-command reproduction pipeline under MIT license.

Keywords: LLM evaluation, peer evaluation, multi-judge, model comparison, benchmark contamination, judge disagreement

Code, data, and reproduction pipeline:
<https://github.com/themultivac/multivac-evaluation>
Interactive platform: <https://app.themultivac.com>

1 Introduction

How good is a language model? The question is deceptively simple. In practice, the answer depends on what you ask, how you score the response, and—critically—who does the scoring. The dominant paradigm for LLM evaluation relies on either static benchmarks or single-model judges, both of which introduce systematic distortions that can mislead practitioners, researchers, and the public. This paper proposes and validates an alternative: a blind peer evaluation methodology in which frontier models simultaneously generate responses and judge each other, producing a multi-perspective assessment that no single evaluator can provide.

1.1 The Benchmark Crisis

Static benchmarks have been the primary instrument for comparing language models since the introduction of GLUE (Wang et al., 2019) and its successors. MMLU (Hendrycks et al., 2021), HumanEval (Chen et al., 2021), and GSM8K (Cobbe et al., 2021) remain widely cited as evidence of model capability. Yet these benchmarks suffer from a fundamental vulnerability: as models train on increasingly large and poorly documented corpora, the probability of test set contamination rises monotonically.

The contamination problem has moved from theoretical concern to documented practice. The retirement of the HuggingFace Open LLM Leaderboard in March 2025 was driven in part by widespread suspicion that leading models had been optimized against its test sets. More directly, Singh et al. (2025) documented that Meta tested 27 private Llama-4 variants on Chatbot Arena before public release, publishing only the highest-performing result. Separately, they estimate that privileged access to Arena data can yield relative performance gains of up to 112% on the Arena distribution, demonstrating how data access asymmetries compound selective publication to distort public rankings.

Even absent deliberate gaming, static benchmarks face a ceiling problem. As frontier models approach perfect scores on MMLU and HumanEval, these benchmarks lose discriminative power. LiveBench (White et al., 2025) addresses this by regularly refreshing its question set, but the core limitation remains: any fixed evaluation protocol becomes a target for optimization.

1.2 The Single-Judge Problem

The LLM-as-a-Judge paradigm (Zheng et al., 2023) emerged as a scalable alternative to human evaluation. MT-Bench and Chatbot Arena demonstrated that GPT-4 judgments correlate reasonably with human preferences, establishing LLM-as-a-Judge as a practical evaluation tool.

However, the same study that validated the approach also documented its failure modes. Three biases are well-established:

Self-enhancement bias. Zheng et al. (2023) observe that GPT-4’s win rate for its own outputs is approximately 10 percentage points higher than under human judging, and Claude-v1’s approximately 25 points higher, though they note that the available data are insufficient to definitively establish self-enhancement bias. Subsequent work has confirmed the pattern at larger scale: models systematically rate outputs from their own family more favorably, suggesting they internalize quality criteria correlated with their own generation patterns.

Position bias. Responses presented first in a pairwise comparison are rated higher regardless of quality. Zheng et al. (2023) found this effect to be statistically significant across evaluation categories.

Verbosity bias. Longer responses receive higher scores independent of content quality. This creates a perverse incentive: models optimized for evaluation performance learn to produce verbose outputs that score well on surface metrics while adding little substantive value. The clarity–depth gap we document

in §5.2—where models score up to 1.8 points higher on clarity than on depth—may partially reflect this dynamic.

The Panel of LLM Judges (PoLL) approach (Verga et al., 2024) demonstrated that using multiple judge models reduces these biases and correlates better with human judgments than any single judge. Our work extends this insight into a fully symmetric framework where every model serves as both contestant and judge.

1.3 Our Contribution

We make three contributions:

1. The Blind Peer Matrix methodology. We formalize an $N \times N$ multi-judge evaluation protocol in which N frontier models generate responses to a shared prompt and then evaluate each other’s outputs under blinded conditions. Self-judgments are excluded. For a pool of $N = 10$ models, each evaluation produces 90 judgment slots, each scored across five weighted dimensions. The methodology is fully specified in §3 and implemented as open-source software.

2. A confound-corrected estimator for same-vendor judge bias. We apply the Blind Peer Matrix across 286 evaluations spanning 198 unique questions in 9 category pools, producing 22,252 valid scored judgments from 55 model configurations. We show that the standard estimator for same-vendor rating bias—the gap between a judge family’s mean score to its own siblings and to others—is confounded by respondent quality and judge leniency, and we decompose it algebraically into these components. Re-estimating with a within-response fixed-effects model with standard errors clustered two-way by judge and question, same-vendor favoritism that survives the controls holds for only two of eight testable families—Anthropic (+0.41, robust to dropping any single sibling judge, worst-case leave-one-judge-out $p < 10^{-4}$) and MiniMax (+0.40, significant at $p = 0.001$ but judge-dependent: its two-way significance does not survive leave-one-judge-out, worst-case $p \approx 0.02$); Qwen (+0.56) does not survive the two-way clustering. The large naive negatives (Mistral -1.02 , Google -0.59) are artifacts, not favoritism.

3. An open disagreement dataset. We release the complete evaluation dataset—27,540 judgment slots (23,356 parsed, of which 22,252 usable-scored; 2,781 intentional self-exclusions; 1,403 judge failures), all question prompts, model responses, per-dimension scores, and judge identities—under MIT license. To our knowledge, this is the largest publicly available multi-judge LLM evaluation dataset with full provenance.

2 Related Work

We situate our work at the intersection of three lines of research: LLM evaluation frameworks, multi-judge approaches, and the emerging study of systematic disagreement between models.

2.1 LLM Evaluation Frameworks

Static benchmarks. The dominant paradigm evaluates models against fixed test sets. MMLU (Hendrycks et al., 2021) tests factual knowledge across 57 domains; HumanEval (Chen et al., 2021) and MBPP (Austin et al., 2021) evaluate code generation; GSM8K (Cobbe et al., 2021) targets mathematical reasoning. HELM (Liang et al., 2023) aggregated multiple benchmarks into a unified evaluation harness. These benchmarks enabled rapid progress but share a structural vulnerability: once published, they become optimization targets.

LLM-as-a-Judge. Zheng et al. (2023) introduced MT-Bench and Chatbot Arena, demonstrating over 80% agreement with human preferences. G-Eval (Liu et al., 2023) extended the approach with chain-of-thought scoring. Prometheus 2 (Kim et al., 2024) trained a dedicated evaluation model to follow rubrics.

Dynamic benchmarks. LiveBench (White et al., 2025) addresses contamination by regularly refreshing its question set. Arena-Hard (Li et al., 2024b) curates difficult questions from Chatbot Arena. Both retain single-judge scoring, inheriting the biases documented in §1.2.

Leaderboard gaming. Singh et al. (2025) provided direct evidence that evaluation gaming extends beyond contamination, documenting selective publication from privately tested model pools.

2.2 Multi-Judge Approaches

Verga et al. (2024) introduced the Panel of LLM Judges (PoLL), demonstrating that a panel of diverse models reduces intra-model bias and correlates better with human judgments than any single judge. Our Blind Peer Matrix extends PoLL in two directions: full symmetry (every model is both contestant and judge) and scale (90 judgment slots per evaluation from 10 simultaneous judges).

JudgeBench (Tan et al., 2024) creates a benchmark for evaluating evaluator models. Two comprehensive surveys (Li et al., 2024a; Gu et al., 2024) catalogue the growing LLM-as-a-Judge literature, identifying single-judge bias as the central unresolved problem.

2.3 Disagreement and Uncertainty

A critical question for any multi-judge system is whether inter-judge disagreement is noise or signal. Shi et al. (2024) provided evidence for the latter in a large-scale study of 15 LLM judges across over 150,000 evaluation instances. They found full unanimity in only approximately 23% of cases, with models sharing architecture and training lineage exhibiting systematically higher internal agreement—a family clustering effect that suggests disagreement reflects genuine differences in evaluative criteria rather than random noise.

Our disagreement analysis (§5.4) extends this by documenting that disagreement varies systematically by task category. The broader uncertainty quantification literature (Gal and Ghahramani, 2016; Xiao and Wang, 2019) has explored single-model confidence, but our peer matrix provides an external, multi-perspective measure that does not depend on any individual model’s calibration.

2.4 Positioning

Three properties distinguish our approach: (1) full symmetry—every model is both judge and contestant; (2) scale—90 judgment slots per evaluation from 10 simultaneous judges; and (3) transparency—we release all judgments, scores, and prompts.

3 Methodology

3.1 Formal Definition

Let $M = \{m_1, m_2, \dots, m_N\}$ be a set of N language models and let q be a question prompt drawn from a category-specific question bank Q .

Response Generation. Each model $m_i \in M$ generates a response $r_i = m_i(q)$. All N models receive the identical prompt q with no system-prompt variation.

Judgment. Each model $m_j \in M$ evaluates each response r_i where $i \neq j$, producing a score vector:

$$\mathbf{s}_{ij} = m_j(r_i, C) = (\text{correctness}_{ij}, \text{completeness}_{ij}, \text{clarity}_{ij}, \text{depth}_{ij}, \text{usefulness}_{ij})$$

The constraint $i \neq j$ enforces **self-exclusion**. This produces an $N \times N$ score matrix S where entry $S_{ij} = \mathbf{w} \cdot \mathbf{s}_{ij}$ is the weighted composite score, and the diagonal is empty. For $N = 10$, each evaluation produces 90 judgment slots.

Aggregation. The final score for respondent m_i on question q is:

$$\text{score}(m_i, q) = \frac{1}{N-1} \sum_{j \neq i} S_{ij}$$

3.2 Scoring Rubric

Judges evaluate each response along five dimensions using a 0–10 scale. The composite weighted score is:

$$S_{ij} = 0.25 \cdot \text{correctness}_{ij} + 0.20 \cdot \text{completeness}_{ij} + 0.20 \cdot \text{clarity}_{ij} + 0.20 \cdot \text{depth}_{ij} + 0.15 \cdot \text{usefulness}_{ij}$$

Correctness receives the highest weight; usefulness the lowest due to susceptibility to verbosity bias. The rubric is provided to judges as a structured system prompt specifying each dimension’s definition and weight. The exact judge prompt is reproduced in Appendix D.

3.3 Blinding and Anti-Contamination

Blinding protocol. Three measures remove explicit identity cues, but they cannot eliminate *stylistic* identification—a judge may still recognize the characteristic style of a model, or of its own family, from the response text alone: (1) *Identity masking*—judges receive only raw response text, stripped of model names, vendor labels, and metadata; (2) *Order randomization*—response order is randomized per judge per evaluation; (3) *Low-temperature judging*—response generation uses temperature 0.7 (default) to preserve natural diversity, while judge API calls use temperature 0.3 to reduce stochastic variation in scoring while avoiding edge-case behavior some providers exhibit at temperature 0. That residual stylistic identification does occur is precisely what our corrected same-vendor analysis detects (§5.5): with explicit cues removed, two families still score their own siblings’ responses above others’ under the within-response fixed-effects estimate (Anthropic +0.41, robust; MiniMax +0.40, significant but judge-dependent; §5.5), evidence that judges partially identify sibling outputs from style.

Anti-contamination. All 198 questions are written de novo (no existing benchmarks) and designed to require cross-domain synthesis. Because they are authored fresh rather than drawn from public benchmarks, training-set contamination is substantially reduced—but not eliminated, for two reasons. First, several providers train on API traffic, so questions submitted through their endpoints may enter future training corpora. Second, all 198 questions are released publicly under MIT with this paper, which makes contamination of *future* model generations a certainty rather than a risk; the dataset is a snapshot whose discriminative power will decay, and re-use of these questions to evaluate later models should not be assumed valid. A subset (66 of the 198) were re-evaluated across different model pools or on different dates, which does not itself reintroduce contamination. The question bank grew from 60 to 198 questions over the evaluation period.

3.4 Category-Specific Model Pools

We compose category-specific pools of $N = 10$ models: six primary categories (code, reasoning, analysis, communication, meta-alignment, edge cases) and three focused batches (SLM, Qwen, MiniMax). Pool composition is detailed in §4.1.

3.5 Head-to-Head Protocol

Separately from the peer matrix, we also ran head-to-head evaluation with a single external judge (Claude Opus 4.6) for rapid pairwise comparison. Because it is single-judge, it does not validate the peer matrix; H2H results are reported (§5.6) as a preliminary side experiment only.

3.6 Implementation

The evaluation framework is implemented in Python, orchestrating API calls via OpenRouter and direct provider APIs. Results are stored as structured JSON with full provenance. The complete pipeline, dataset (27,540 judgments), all 198 question prompts, and scoring rubric are released under MIT license.¹

¹Code and dataset: <https://github.com/themultivac/multivac-evaluation>. Platform: <https://app.themultivac.com>.

4 Experimental Setup

4.1 Models

A total of 55 model configurations (50 distinct by display name) participated. The core pool of 10 frontier models (GPT-5.4, Claude Opus 4.6, Claude Sonnet 4.6, GPT-OSS-120B, MiMo-V2-Flash, Grok 4.20, DeepSeek V3, Gemini 3.1 Pro, Gemini 3 Flash Preview, MiniMax M2.5) appears in 163–238 evaluations each. An extended pool of 11 models participates in 30–60 evaluations, and 34 focused batch models appear in 1–25 evaluations. The full model list is in Appendix A.

Rankability. For any *global* ranking we require a model to have participated in ≥ 3 category pools and to have ≥ 30 within-judge comparisons. Bradley–Terry estimates a single latent ability; a single-category model would confound that ability with its category’s difficulty and judge composition, so single-category models are ranked within their category only (this rule is justified by the estimator’s assumption, not by the ranking it produces). Collapsing the 55 registry keys to 50 distinct display names and applying this rule yields a 30-model frontier pool and an 18-model broad-participation generalist pool; the two configurations with no scored judgments (Qwen 3 235B-A22B, Qwen3-Next 80B-A3B) are unrankable, leaving 48 of the 50 distinct models placed. (The “34 focused batch models” above is a participation tier, distinct from the 30-model frontier ranking pool.)

4.2 Question Design

All 198 questions were authored *de novo* following four principles: discriminative power (targeting the capability frontier), cross-domain synthesis (reducing memorization advantage), category alignment (one category per question), and fresh authorship (no question drawn from an existing benchmark). Distribution of the 198 unique questions: code (30), analysis (30), communication (30), reasoning (30), meta-alignment (30), SLM (14), MiniMax (13), Qwen (11), edge cases (10). These 198 questions were run in 286 peer-matrix evaluations: the five primary pools each hold 30 questions but were evaluated 44–50 times each, because 66 questions were re-evaluated (with different model pools or on different dates, up to six times each), while the four focused batches were evaluated once per question. Sample questions are in Appendix B.

4.3 Evaluation Protocol

Evaluations were conducted February–April 2026 across two phases. Phase 1 (February–March) iteratively refined the scoring rubric, judgment prompt, and parsing logic; Phase 2 (March–April) applied the mature protocol at scale. Two Phase 1 judgments containing scores of 100 (on a 0–10 scale) due to absent range clamping were identified and excluded. Over the frozen dataset, the model-attributable parse-failure rate is 4.9% of attempted non-self judgments (1,192 parse failures over the 24,548 non-self judgments that reached a model; the 211 API-side failures folded into the 1,403 total of §4.4 are excluded from both numerator and denominator, so this is not 1,403/24,759), and it is strongly systematic across both judge models ($\chi^2(49) = 7,122$, $p \approx 0$) and categories ($\chi^2(8) = 673.1$, $p = 4.3 \times 10^{-140}$): a few judges fail most of the time (olmo_think 85.6%, gemini_3_pro 56.0%) while others never fail, and per-category failure rates range from 19.2% (Qwen, the highest) and 12.0% (edge cases) down to 1.6% (communication) and 1.4% (MiniMax, the lowest), with the five primary pools falling between—7.2% (reasoning), 4.9% (code), 4.2% (analysis), 1.9% (meta-alignment), 1.6% (communication)—and the small focused pools (Qwen, edge cases, SLM 9.5%, MiniMax) at the extremes on fewer than 1,000 judgments each. See Appendix D for the full Phase 1/Phase 2 protocol comparison.

Cost. Each peer matrix evaluation costs \$2–3 in API fees. The complete dataset cost approximately \$700–850.

4.4 Judgment Parsing

The parsing pipeline extracts per-dimension scores using a multi-strategy approach: thinking-block removal, JSON brace matching, regex fallback, and 0–10 range clamping. Every reported number derives from an explicit, reproducible filter over the 27,540 raw judgment slots:

Total judgment slots	27,540
– intentional self-exclusions (matrix diagonal)	–2,781
– judge failures (parse/API/clamp)	–1,403
Successfully parsed judgments	23,356
– valid zero-scores (refused/empty/broken responses)	–1,104
Usable scored judgments	22,252
– unmapped-vendor judgments	–839
Same-vendor analysis set (§5.5)	21,413

The 1,104 valid zero-scores are genuine judge ratings of non-responses (the parser returns a score only when at least three of the five dimensions are explicitly present; parse failures set an error flag). They are included in the §5.1 leaderboard—where a refusal is a legitimate zero and we report per-model refusal rates—but excluded from the §5.5 same-vendor analysis, where an all-zero response carries no within-response information. The complement of the previously reported “22,254 valid / 5,286 self-excluded” conflated 2,781 self-exclusions with 1,403 failures and 1,104 zero-scores; the true self-exclusion count is 2,781.

5 Results

We report results across 286 peer matrix evaluations spanning 198 unique questions and 9 category pools, yielding 22,252 valid scored judgments from 55 model configurations. Intentional self-judgments ($N = 2,781$) were excluded per protocol; a further 1,403 judgments were genuine judge failures and 1,104 were valid zero-scores of refused responses (§4.4).

5.1 Overall Cross-Category Rankings

Rankings depend heavily on how judgments are aggregated, and two corrections to the naive approach change the picture materially. First, **genuine zero-scores are included**. The 1,104 judgments in which a judge scored a refused, empty, or truncated response zero are real ratings, not parse failures (§4.4); excluding them inflates the mean of models that frequently fail to answer—by up to four points (rates and means here are computed over each model’s complete cross-pool judgment set: MiniMax M2.1, 53% refusal rate: naive mean 7.53 $\rightarrow$ 3.51; Qwen 3 32B, 42%: 9.13 $\rightarrow$ 5.33; Kimi K2.5, 43%: 8.63 $\rightarrow$ 4.94). We include zero-scores and report per-model refusal rates. Second, **aggregate mean scores are confounded by judge leniency** (§5.5); we therefore report a leniency-adjusted Bradley–Terry (BT) ranking alongside the naive mean, computed from within-judge pairwise comparisons so that judge leniency and item difficulty cancel by construction.

No category has a separated winner. We rank each pool with a leniency-adjusted Bradley–Terry (BT) model on within-judge pairwise comparisons, so judge leniency and item difficulty cancel by construction, and we bootstrap at the question level (66 of the 198 questions recur; §4.2). Table 1 gives each pool’s top model against its runner-up: in all nine pools the winner beats second place in only 50–58% of matched comparisons, and the 95% interval on that win probability spans 50% in every case—not one pool has a statistically separated winner. Four different models lead the six primary categories by point estimate (GPT-5.4 in analysis, code, and reasoning; Mistral Small Creative in communication; Gemini 3 Flash in edge cases; Claude Opus 4.6 in meta-alignment), but none separates from second. This is both because the top models are close and because ~ 30 questions per pool cannot resolve gaps of this size (§6.3); we report it as an unresolved comparison, not a demonstrated tie.

The champion is an artifact of aggregation. Aggregating across categories does separate a winner, but which winner depends entirely on the aggregation (Figure 1). By naive mean the top models are GPT-OSS-120B-Legal, Grok 4.1 Fast, and Grok Code Fast 1—scored disproportionately by lenient judges. A leniency-adjusted global BT ranking over the frontier pool instead crowns a near coin-flip that mixes GPT-5.4 with single-category specialists (Mistral Small Creative, communication-only; Grok Code Fast 1, code-only) whose global strength is category exposure, not ability. Restricting to broad-participation generalists

Table 1: Per-category Bradley–Terry winner vs. runner-up, question-clustered bootstrap. “Win vs. 2nd” is the top model’s pairwise win probability over second place; the difference CI is the 95% interval on that win probability. No pool has a separated winner—every interval spans 50%.

Pool	Winner (2nd)	Win vs. 2nd	95% CI on win	Separated?
code	GPT-5.4 (Grok Code Fast 1)	53%	[38%, 65%]	no
reasoning	GPT-5.4 (Claude Opus 4.6)	53%	[42%, 63%]	no
analysis	GPT-5.4 (Grok 4.20)	51%	[44%, 58%]	no
communication	Mistral Small Creative (Claude Sonnet 4.6)	51%	[42%, 61%]	no
meta-alignment	Claude Opus 4.6 (GPT-5.4)	52%	[46%, 58%]	no
edge cases	Gemini 3 Flash Preview (Claude Opus 4.5)	50%	[37%, 62%]	no
MiniMax	GPT-5.4 (Claude Sonnet 4.6)	53%	[45%, 60%]	no
Qwen	Qwen 3.5 397B-A17B (Qwen 3.5 122B-A10B)	53%	[34%, 73%]	no
SLM	Qwen 3 8B (Gemma 3 27B)	58%	[48%, 71%]	no

Separated winners: 0 of 9 pools

(≥ 3 categories, §4.1) removes that confound: GPT-5.4 holds rank 1 in 99% of question-clustered resamples and is the sole rank-1 coverage member, unchanged when its one extra category is dropped and when the comparison is restricted to the five categories its rivals share. Its per-comparison edge over the runner-up (Grok 4.20) is a modest 54% win rate (CI 51–58%), and Grok 4.20 is not excludable under family-wise correction—so GPT-5.4 leads the generalists but does not dominate them.




Figure 1: The aggregation ladder: the frontier models ranked three ways—naive mean, leniency-adjusted Bradley–Terry, and exposure-restricted BT (generalists with ≥ 3 categories). GPT-5.4 (green) is mid-pack by naive mean but rank 1 under both BT variants; single-category specialists (orange), high under global BT, drop out of the exposure-restricted column because they enter only one category. The champion changes at every rung, which is the point: the “overall” winner is an artifact of the aggregation. Generated by `make_finding1.py` from the frozen data.

No single global leaderboard exists. The comparison graph is not fully connected: after collapsing duplicate registry keys to distinct models (§4.1), the 30-model frontier pool and the 18-model small-model batch pool share no cross-pool comparisons, so no single 48-model leaderboard is identifiable; rankings are reported within pool and per category. Even in the frontier pool, where a global ranking *is* identifiable, it is exposure-confounded (the aggregation ladder above)—which is why the per-category rankings, not any global aggregate, are the primary result.

Dropping non-responses is a confounded practice. This is a general methodological point, not a per-model curiosity. Standard leaderboard construction silently discards judgments that failed to yield a score—but genuine ratings of refused or empty responses are among them (§4.4), and discarding them systematically favors models that refuse or fail to produce parseable output. It is structurally the same failure as the same-vendor estimator (§5.5): a routine aggregation confounded by a nuisance factor, demonstrated by counterexample. Table 2 and Figure 2 quantify the effect. The frontier top tier (near-zero refusal) is unaffected, while high-refusal models fall by up to four points once their zero-scores are counted.

Table 2: Effect of counting non-responses on the mean composite score, for the highest-refusal models. “Scored-only” is the mean over judgments with a positive score (the previously standard aggregation); “incl. non-responses” adds the genuine zero-scores of refused/empty responses. The frontier top tier (GPT-5.4, both Claudes, Grok, Mistral Small; 0% refusal) is unchanged. Refusal rate and both means are computed over all judgments each model received across pools; the SLM-pool-restricted figures in §5.2 differ for models judged in more than one pool (e.g. Qwen 3 32B, 39% refusal within the SLM pool vs. 42% across all pools).

Model	Refusal rate	Scored-only	Incl. non-responses
MiniMax M2.1	53%	7.53	3.51
Kimi K2.5	43%	8.63	4.94
MiniMax M2	42%	7.21	4.21
Qwen 3 32B	42%	9.13	5.33
Olmo 3.1 32B Think	37%	7.94	4.99




Figure 2: Effect of counting non-responses, per model: scored-only mean (grey) vs. mean including the genuine zero-scores of refused/empty responses (green); each model’s refusal rate is printed above its bars. The five highest-refusal models drop by up to four points, while the zero-refusal frontier tier (right) is unchanged—a flat control. Same quantities as Table 2; generated by `make_fig_refusal.py` from `finding1_leaderboard.json`.

5.2 Category-Specific Analysis

Under the leniency-adjusted Bradley–Terry ranking, four distinct models lead the six primary category pools by point estimate—GPT-5.4 in analysis, code, and reasoning; Mistral Small Creative in communication; Claude Opus 4.6 in meta-alignment; Gemini 3 Flash in edge cases—though none separates from its runner-up (Table 1). Of the three focused batches only the Qwen pool is single-family and led within family (Qwen 3.5 397B); the SLM pool is a multi-family field of eight vendors led by Qwen 3 8B, and the MiniMax batch pairs MiniMax models against frontier judges and is led by GPT-5.4. The naive mean produces a different winner in most pools, illustrating how judge leniency reshapes per-category rankings, not only the aggregate.

Small Language Models. Within the SLM pool ($\leq 32\text{B}$ parameters, zero-scores included), Qwen 3 8B (9.343 within the SLM pool; 9.15 across all pools, Table 4) outperforms models with $2\text{--}4\times$ its parameter count, including Phi-4 14B (8.917) and Gemma 3 27B (8.828). This ordering is by raw within-pool mean and is not leniency-adjusted (§5.5); under the within-pool Bradley–Terry ranking the headline result holds—Qwen 3 8B is BT rank 1 of the small-model pool—but the raw-mean order of the mid-pool models is not leniency-robust, as BT places Gemma 3 27B above Phi-4 14B. Including refused responses sharpens the headline gap: the larger Qwen 3 32B falls from 8.934 to 5.438 (39% refusal within the SLM pool; its all-pool refusal rate is 42%, Table 2) and Kimi K2.5 from 8.628 to 4.940 (43%), so the 8B model outscores its 32B sibling by nearly four points once non-responses are counted.

Dimension-level analysis. Across all judgments (zero-scores included), depth is consistently the lowest-scoring dimension (mean: 7.622) while clarity is the highest (mean: 8.583). The clarity–depth gap exceeds 0.8 points for 25 of the 33 models with ≥ 100 judgments, with Llama 3.1 8B showing the largest gap ($\Delta = 1.780$), ahead of Mistral Nemo 12B (1.505) and GPT-5.2-Codex (1.504). This suggests current models are better calibrated for well-structured output than for deep analytical engagement.

5.3 Judge Behavior Analysis

Judge behavior varies dramatically. The strictest high-volume judge is GPT-5.4 (mean 7.187 across 1,565 scored judgments); the most lenient among the frontier judges is Mistral Small Creative (9.695 across 422 scored judgments)—a spread of 2.51 points. We note for accuracy that Mistral Small is the most lenient only within this curated frontier set; across all 48 judges with ≥ 30 judgments, Gemini 2.5 Flash-Lite is more lenient still (9.765). OpenAI models exhibit high variance as judges (GPT-5.4 std = 2.215), while Seed 1.6 Flash (std = 0.877) and Granite 4.0 Micro (std = 0.488) compress scores into a narrow band.

5.4 Disagreement Patterns

Code evaluation produces the highest average disagreement (mean within-response $\sigma = 1.269$), just over double that of meta-alignment ($\sigma = 0.632$; ratio $2.01\times$). Edge cases show high mean disagreement (1.033) but a notably lower median (0.665), indicating a right-skewed distribution with occasional extreme disagreements. This mean–median gap holds across all categories: most evaluations produce moderate agreement, but a tail of high-disagreement cases exists in every category (Figure 3); overall inter-annotator agreement is quantified in §6.6.

5.5 Same-Vendor Rating Bias

The naive estimator is confounded. The standard approach estimates same-vendor bias as the gap between the mean score a judge family gives its own siblings and the mean it gives other vendors. Writing the four relevant cells as A (same-vendor judge scoring a same-vendor respondent), B (same-vendor judge, other respondent), C (other judge, same-vendor respondent), and D (other judge, other respondent), the naive estimate $A - B$ decomposes exactly as

$$A - B = \underbrace{(A - C)}_{\text{favoritism}} + \underbrace{(C - D)}_{\text{respondent quality}} - \underbrace{(B - D)}_{\text{judge leniency}} .$$




Figure 3: Within-response disagreement σ by category: mean (bars) and median (red markers). Code shows the highest mean disagreement (1.269), just over double meta-alignment’s (0.632); the mean–median gap—largest for edge cases (1.033 vs. 0.665)—reflects a right-skewed tail of high-disagreement evaluations. Generated by `make_fig_disagreement.py` from `category_disagreement.json`.

The respondent-quality term ($C - D$) measures how good a family’s answers are, as judged by neutral judges; the leniency term ($B - D$) measures how generous its judges are toward everyone. Neither is favoritism. We verify this identity holds to 10^{-9} across all eight families and confirm our $A - B$ values reproduce the previously published naive estimates exactly.

Defining the estimand. Our reported estimate is the coefficient on the same-vendor indicator in a within-response fixed-effects model that absorbs the response and adjusts for judge-level baseline strictness, with standard errors clustered two-way by judge and question—the primary specification throughout (Table 3). Where judge-only clustering would change a verdict, we say so explicitly. This isolates a judge’s differential treatment of a sibling’s output, holding both the output and judge strictness fixed. The distinction matters: for OpenAI the raw $A - C$ gap is -1.136 , but OpenAI judges are the strictest in the pool (leniency -1.052), so the corrected estimate is -0.137 and not significant.

Two mechanisms produce large but spurious naive estimates. **Google** (-0.593 naive) is a respondent-quality artifact: neutral judges score Google’s responses 0.938 below others ($C - D$), so Google judges score them lower simply because they are weaker; corrected, $+0.209$ ($p = 0.058$, not significant). **Mistral** (-1.017 naive, the largest absolute value) is a judge-lenience artifact: Mistral judges score all respondents 1.094 above the neutral baseline ($B - D$), and Mistral’s own responses are above average ($C - D = +0.527$); corrected, -0.164 ($p = 0.345$)—the apparent bias does not reverse under controls, it collapses toward zero. **xAI** ($+0.745$ naive) is a false positive driven by strong responses, corrected -0.095 ($p = 0.549$).

Corrected estimates. Under the primary specification—standard errors clustered two-way by judge and question, since 66 questions recur (§4.2)—only two families show same-vendor favoritism that survives the controls (Table 3, Figure 4): Anthropic ($+0.406$, $p < 0.001$, 482 identifying responses) and MiniMax ($+0.397$, $p = 0.001$, 50). The two positives are not equally robust. Anthropic’s estimate survives dropping any single sibling judge—its worst-case leave-one-judge-out p stays below 10^{-4} , on 482 identifying responses spread across four Claude judges. MiniMax’s two-way significance is judge-dependent: it rests on 50 identifying responses in a MiniMax-heavy pool where six sibling judges overlap on the same responses, and removing individual MiniMax judges raises its worst-case two-way p to ≈ 0.02 (dropping the MiniMax M2.5 judge), just past the Bonferroni threshold. We therefore report MiniMax as a positive with this stated fragility, not as robust on par with Anthropic. Qwen ($+0.557$) is significant under judge-clustering alone ($p = 0.002$) but

Table 3: Same-vendor bias: naive estimate ($A - B$) vs. corrected within-response estimate. Bold = significant under the within-response fixed-effects model with standard errors clustered *two-way by judge and question* (the primary specification, since 66 questions are re-evaluated and their judgments correlate) and Bonferroni correction ($\alpha = 0.00625$). n_{id} is the number of responses judged by both a sibling and a non-sibling. Only Anthropic and MiniMax survive; Qwen ($p = 0.018$) does not. Intervals are uncorrected 95% CIs on the corrected estimate under the same two-way specification; because significance is judged at the Bonferroni level $\alpha = 0.00625$, an interval may exclude zero while the family is still not counted (Qwen). The full four-cell decomposition ($A/B/C/D$) is in Appendix F.

Family	Naive ($A - B$)	Corrected (FE)	95% CI (two-way)	p (judge+question)	n_{id}
Anthropic	+0.616	+0.406	[+0.312, +0.500]	< 0.001	482
MiniMax	+0.314	+0.397	[+0.154, +0.641]	0.001	50
Qwen	+0.913	+0.557	[+0.096, +1.017]	0.018	22
Google	-0.593	+0.209	[-0.017, +0.435]	0.070	394
Meta	-0.681	+0.195	[-0.081, +0.472]	0.166	26
OpenAI	+0.229	-0.137	[-0.404, +0.130]	0.316	385
xAI	+0.745	-0.095	[-0.409, +0.219]	0.554	59
Mistral	-1.017	-0.164	[-0.497, +0.170]	0.336	25




Figure 4: Naive ($A - B$, open markers) vs. corrected within-response fixed-effects (filled) same-vendor estimate per family. Green = significant under the primary specification (two-way judge+question clustering, Bonferroni)—the same two families, Anthropic and MiniMax, as Table 3. Four families cross zero under correction (Google, xAI, OpenAI, Meta); Mistral collapses toward zero without crossing. Generated from the released data by `make_fig_bias.py`.

not once question-level correlation is admitted ($p = 0.018$), and it rests on 22 identifying responses from a focused Qwen-family pool with a negative raw $A - C$ gap; we do not count it. Google, OpenAI, xAI, Meta, and Mistral show no significant effect. The naive estimator called five families positive and three

negative; under controls, one is robustly positive (Anthropic), one is positive with stated judge-dependence (MiniMax), and none are negative.

What the estimator cannot separate. We name the interpretation this design cannot rule out. After the response is absorbed as a fixed effect and judge leniency is adjusted for, sibling judges still score sibling outputs higher. That residual is consistent with *favoritism*—a judge treating a sibling’s output preferentially—but it is equally consistent with *stylistic affinity that tracks genuine quality*: models drawn from a shared training distribution may parse a sibling’s formatting, notation, and argument structure more accurately, and so score it more correctly rather than more generously. Our estimator separates favoritism from respondent quality and from judge leniency, which is what the four-cell decomposition establishes; it does not separate favoritism from same-distribution comprehension. Distinguishing the two requires an external anchor—human judgments on the same responses, against which sibling and non-sibling scores could be checked for differential accuracy rather than differential level. We do not have that anchor here, and §6.1 flags it as the necessary next study. The corrected estimates should therefore be read as an upper bound on same-vendor favoritism, not as a point measurement of it.

Robustness and power. Excluding the five judges that fail to produce parseable output more than 30% of the time and re-fitting leaves every family verdict unchanged. The per-category rankings (§5.1) are likewise robust to the same-vendor effect: dropping all same-vendor judgments and refitting each pool leaves every pool’s point-estimate leader unchanged (the single-family Qwen batch has no cross-vendor judgments and is excluded from this check). xAI (59 identifying responses), Meta (26), Mistral (25), and Qwen (22) have small samples; their null results should be read as an absence of demonstrated bias, not demonstrated absence of bias. Meta’s corrected estimate (+0.195, $p = 0.108$) is positive and directionally suggestive but underpowered.

5.6 Head-to-Head Side Experiment

As a preliminary side experiment—not part of the peer-matrix results, and not a validation of them—we ran head-to-head batches with Claude Opus 4.6 as a single judge (185 recorded questions; two batches, of 151 and 31 questions, carry substantive volume). The largest batch paired Qwen 3.6 Plus against DeepSeek V3.2 over 151 questions, with Qwen winning 107–11 and 26 ties (7 incomplete). Two limits bound what this can show. First, it is *single-judge*, and is therefore subject to precisely the leniency and idiosyncrasy effects documented in §5.3—the failure mode the peer matrix exists to avoid—so it cannot corroborate the peer-matrix rankings. Second, Qwen 3.6 Plus is not part of the peer-matrix registry: it appears in no peer-matrix evaluation and is absent from Table 4, so this comparison does not involve the study’s model set. We report the batches only to document the released head-to-head files.

6 Limitations

6.1 LLMs Evaluating LLMs

Multi-judge aggregation reduces individual biases but does not eliminate biases shared across all judges. The clarity–depth gap (§5.2) may reflect exactly this kind of shared bias. A human correlation study is needed—and is also what would be required to separate same-vendor favoritism from same-distribution comprehension in §5.5, where the two are observationally equivalent under our design.

6.2 Output Quality vs. Reasoning Process

The methodology evaluates final output quality without examining the reasoning process. A model that arrives at a correct answer through flawed reasoning receives the same score as one demonstrating genuine understanding. Future work should incorporate reasoning trace evaluation (Turpin et al., 2023).

6.3 Sample Size and Statistical Power

Per-category and per-model sample sizes vary considerably. Core pool models participate in 163–238 evaluations; focused batch models may appear in as few as one. Bootstrap confidence intervals (question-clustered, 10,000 resamples) quantify ranking uncertainty within the 30-model frontier pool (§5.1); 21 models participate in ≥ 30 evaluations each.

The per-category rankings are underpowered, and we are explicit about what that does and does not license. With ~ 30 questions per pool, the minimum detectable difference between adjacent models is a pairwise win rate of roughly 59–76% (0.35–1.17 log-ability units), whereas the observed top-two gaps are only 50–58% win rates—2 to 73 times below the detection floor—so we cannot separate the top model from the runner-up in any pool. This does *not* mean the models are equal: the difference intervals are wide enough to admit a genuine advantage in either direction (e.g. up to a 58% win rate for second place in meta-alignment). Reading an underpowered null as evidence of no difference is the same absence-of-evidence error that inflates the naive same-vendor estimator (§5.5) and that we flag for the small-sample families there: an unresolved comparison is not a demonstrated tie. Only aggregation across categories, which pools each model’s comparisons past this floor, separates a winner (GPT-5.4 among generalists, §5.1). That separation is contingent on the exposure restriction we impose: the ≥ 3 -category cutoff defining the generalist pool (§4.1) is an analyst-chosen threshold. The GPT-5.4 result is unchanged when its one extra category is dropped (§5.1), but the cutoff itself is a researcher degree of freedom the reader should weigh.

6.4 Question Design Bias

The first author wrote all 198 questions, introducing unavoidable perspective bias. We mitigate through category diversity and volume; questions were authored fresh (not drawn from existing benchmarks) over the evaluation period, though a subset were re-evaluated across pools (§4.2). Releasing the question bank is itself a limitation: publishing all 198 questions under MIT guarantees that future models may be trained on them, so this benchmark’s discriminative power decays from the moment of release (§3.3). The frozen results reported here are valid for the models and dates evaluated; re-running these questions against later model generations should not be assumed to measure the same thing.

6.5 Cost and Throughput

At \$2–3 per evaluation, the methodology is expensive relative to static benchmarks but inexpensive relative to human evaluation.

6.6 Inter-Annotator Agreement

We compute Krippendorff’s α (interval level) across all judge–respondent pairs. The overall $\alpha = 0.618$, approaching the conventional threshold of 0.667 for tentative conclusions (Krippendorff, 2011). Two categories—reasoning ($\alpha = 0.687$) and MiniMax ($\alpha = 0.673$)—exceed this threshold, while communication ($\alpha = 0.354$) and edge cases ($\alpha = 0.436$) show lower agreement. The reasoning result is notable: judges converge more readily on logical analysis quality than on communication quality, consistent with the intuition that reasoning has more objective markers.

The moderate overall agreement supports our core methodological argument: single-judge evaluation is unreliable precisely because judges disagree systematically. An α of 0.618 indicates judges share enough common ground to produce meaningful aggregate rankings while disagreeing enough that no single judge’s perspective dominates. Both α and the category disagreement statistics (§5.4) are computed per judgment and are not corrected for the 66 repeated questions (§4.2), so their nominal precision is somewhat overstated.

7 Discussion

7.1 Implications for Model Selection

Model selection should be task-specific. The finding that GPT-5.4 leads the exposure-restricted generalist ranking yet is not statistically separated from the runner-up in any single category illustrates that raw mean scores, leniency-adjusted rankings, and per-category rankings answer different questions. Category-specific rankings (§5.2) provide more actionable signal than any aggregate.

7.2 Implications for AI Safety

The same-vendor bias results (§5.5) imply that single-family evaluation of alignment properties may produce optimistic results, though the effect is smaller than the naive estimate suggests—at most about +0.4 composite points, robust for one family (Anthropic) and, for the second (MiniMax), significant under the primary specification but judge-dependent under leave-one-judge-out (§5.5). The low disagreement on meta-alignment ($\sigma = 0.632$) may reflect shared RLHF biases rather than genuine consensus.

7.3 The Data Engine

Each evaluation generates 90 scored judgments—convertible to preference pairs—usable for DPO/RLHF training at <\$0.01/sample. The five-dimensional scoring enables optimization for specific dimensions rather than monolithic preference.

7.4 Toward Open Evaluation Infrastructure

Independent, reproducible evaluation infrastructure is increasingly needed as models are deployed in higher-stakes settings. We release the evaluation framework, question bank, scoring rubric, and complete dataset under MIT license. At current API prices the full matrix costs \$2–3 per evaluation; we are exploring open-source judge distillation to reduce inference costs by an estimated 70%, which would make continuous evaluation economically viable at scale.

8 Conclusion

We introduced the Blind Peer Matrix, a multi-judge evaluation methodology in which N frontier language models simultaneously generate responses and evaluate each other’s outputs under blinded conditions. Applied across 286 evaluations spanning 198 unique questions and 9 category pools, the method produced 22,252 valid scored judgments from 55 model configurations—to our knowledge, the largest publicly available multi-judge LLM evaluation dataset with full provenance.

Three findings stand out. First, no category has a statistically separated winner—in all nine pools the top model beats second place in only 50–58% of matched comparisons—so the “overall” champion is an artifact of the aggregation choice; only a leniency-adjusted and exposure-restricted ranking of broad-participation generalists separates one (GPT-5.4, rank 1 in 99% of question-clustered resamples and the sole rank-1 coverage member, though its runner-up is not excludable under family-wise correction). Second, the standard estimator for same-vendor judge bias is confounded by respondent quality and judge leniency; under a within-response fixed-effects model with standard errors clustered two-way by judge and question, only two of eight families show favoritism that survives the controls—Anthropic (+0.41, robust under leave-one-judge-out, worst-case $p < 10^{-4}$) and MiniMax (+0.40, significant at $p = 0.001$ but judge-dependent, its two-way significance not surviving leave-one-judge-out, worst-case $p \approx 0.02$)—Qwen (+0.56) does not survive the clustering, and no family shows significant negative bias—the naive negatives for Mistral (−1.02) and Google (−0.59) are artifacts of judge leniency and respondent quality, respectively. Third, judge disagreement is category-dependent ($\sigma = 1.27$ for code vs. 0.63 for meta-alignment), with overall inter-annotator agreement of $\alpha = 0.618$ (Krippendorff’s alpha).

We release the complete dataset (27,540 judgments), all 198 question prompts, the evaluation framework source code, and documentation under MIT license.

References

- Austin, J., Odena, A., Nye, M., Bosma, M., Michalewski, H., Dohan, D., Jiang, E., Cai, C., Terry, M., Le, Q., and Sutton, C. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Chen, M., Tworek, J., Jun, H., et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Cobbe, K., Kosaraju, V., Bavarian, M., et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Gal, Y. and Ghahramani, Z. Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. In *ICML*, pp. 1050–1059, 2016.
- Gu, J., Jiang, X., Shi, Z., et al. A survey on LLM-as-a-judge. *arXiv preprint arXiv:2411.15594*, 2024.
- Hendrycks, D., Burns, C., Basart, S., Zou, A., Mazeika, M., Song, D., and Steinhardt, J. Measuring massive multitask language understanding. In *ICLR*, 2021.
- Kim, S., Shin, J., Cho, Y., et al. Prometheus 2: An open source language model specialized in evaluating other language models. In *EMNLP*, 2024.
- Krippendorff, K. Computing Krippendorff’s alpha-reliability. Departmental Papers (ASC), University of Pennsylvania, 2011.
- Li, D., Jiang, B., Huang, L., et al. From generation to judgment: Opportunities and challenges of LLM-as-a-judge. *arXiv preprint arXiv:2411.16594*, 2024a. Accepted at EMNLP 2025.
- Li, T., Chiang, W.-L., Frick, E., et al. From crowdsourced data to high-quality benchmarks: Arena-Hard and BenchBuilder pipeline. *arXiv preprint arXiv:2406.11939*, 2024b.
- Liang, P., Bommasani, R., Lee, T., et al. Holistic evaluation of language models. *Transactions on Machine Learning Research*, 2023.
- Liu, Y., Iter, D., Xu, Y., Wang, S., Xu, R., and Zhu, C. G-Eval: NLG evaluation using GPT-4 with better human alignment. In *EMNLP*, 2023.
- Shi, L., Ma, C., Liang, W., et al. Judging the judges: A systematic study of position bias in LLM-as-a-judge. *arXiv preprint arXiv:2406.07791*, 2024. Accepted at ACL-IJCNLP 2025.
- Singh, S., Nan, Y., Wang, A., et al. The leaderboard illusion. *arXiv preprint arXiv:2504.20879*, 2025. Accepted at NeurIPS 2025 Datasets and Benchmarks Track.
- Tan, S., Zhuang, S., Montgomery, K., et al. JudgeBench: A benchmark for evaluating LLM-based judges. *arXiv preprint arXiv:2410.12784*, 2024. Accepted at ICLR 2025.
- Turpin, M., Michael, J., Perez, E., and Bowman, S. R. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. In *NeurIPS*, 2023.
- Verga, P., Hofstatter, S., Althammer, S., et al. Replacing judges with juries: Evaluating LLM generations with a panel of diverse models. *arXiv preprint arXiv:2404.18796*, 2024.
- Wang, A., Singh, A., Michael, J., Hill, F., Levy, O., and Bowman, S. R. GLUE: A multi-task benchmark and analysis platform for natural language understanding. In *ICLR*, 2019.
- White, C., Dooley, S., Roberts, M., et al. LiveBench: A challenging, contamination-limited LLM benchmark. In *ICLR*, 2025.

Xiao, Y. and Wang, W. Y. Quantifying uncertainties in natural language processing tasks. In *AAAI*, 2019.

Zheng, L., Chiang, W.-L., Sheng, Y., et al. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. In *NeurIPS Datasets and Benchmarks Track*, 2023.

A Complete Model List

Table 4 lists all 55 entries of the released `models_used` registry, with vendor family, evaluation count, judgments issued as a judge, mean received composite score (zero-scores included), and category participation. The 55 registry entries resolve to 50 distinct models by display name: five names recur under a second key—a judge-only configuration or a re-run variant (e.g. two *Seed 1.6 Flash* and two *Gemini 2.5 Flash* keys). They span 17 vendor families: Anthropic, OpenAI, Google, xAI, DeepSeek, Qwen/Alibaba, MiniMax, Mistral, Meta, Xiaomi, ByteDance, Moonshot, Microsoft, IBM, Ai2, Zhipu, and NVIDIA. Parameter counts are recorded in the repository only for the open ≤ 32 B pool and are therefore omitted here for the proprietary majority.

Table 4: Complete model list (55 model configurations from `models_used`, 50 distinct by display name). Mean is the received composite score with genuine zero-scores included; *n/a* marks roster configurations that produced no parseable judgment in the frozen data (issued and received none). Judg. given counts non-self judgments a model issued as judge, zero-scores included, and can therefore exceed the scored-judgment counts underlying the judge-leniency means in §5.3 (e.g. Mistral Small Creative: 423 issued, 422 scored). Parameter counts are omitted: the repository records them only for the open ≤ 32 B pool, not for the proprietary majority.

Family	Model	Evals	Judg. given	Mean	Categories
Ai2	Olmo 3.1 32B Think	10	13	4.99	reas
Anthropic	Claude Opus 4.6	188	1632	8.43	anly, code, comm, meta, reas
Anthropic	Claude Sonnet 4.6	188	1640	8.48	anly, code, comm, meta, reas
Anthropic	Claude Opus 4.5	60	537	9.08	anly, code, comm, edge, meta, reas
Anthropic	Claude Sonnet 4.5	60	538	9.07	anly, code, comm, edge, meta, reas
Anthropic	Claude Sonnet 4.6	13	91	8.71	mmax
ByteDance	Seed 1.6 Flash	38	340	8.45	comm
ByteDance	Seed 1.6 Flash	10	89	9.43	comm
DeepSeek	DeepSeek V3	188	1612	8.20	anly, code, comm, meta, reas
DeepSeek	DeepSeek V3.2	60	530	9.03	anly, code, comm, edge, meta, reas
Google	Gemini 3.1 Pro	188	1612	6.44	anly, code, comm, meta, reas
Google	Gemini 3 Flash Preview	163	1434	8.74	anly, code, edge, meta, reas
Google	Gemini 3 Pro Preview	50	197	6.61	anly, code, edge, meta, reas
Google	Gemini 2.5 Flash	37	313	7.53	reas
Google	Gemini 2.5 Flash	30	270	8.93	anly, comm, reas
Google	Gemma 3 27B	14	124	8.83	slm
Google	Gemini 3.1 Pro	13	0	n/a	mmax
Google	Gemini 2.5 Flash-Lite	10	89	8.54	comm
Google	Gemma 3n 4B	1	7	1.84	slm
IBM	Granite 4.0 Micro	14	123	8.35	slm
Meta	Llama 4 Scout	14	123	8.45	slm
Meta	Llama 3.1 8B	13	117	7.58	slm
Microsoft	Phi-4 14B	13	117	8.92	slm
MiniMax	MiniMax M2.5	163	1375	5.42	anly, code, meta, mmax, reas
MiniMax	MiniMax M2	23	175	4.21	code, mmax

Family	Model	Evals	Judg. given	Mean	Categories
MiniMax	MiniMax-01	13	91	7.99	mmax
MiniMax	MiniMax M1	13	90	5.82	mmax
MiniMax	MiniMax M2.1	13	88	3.51	mmax
MiniMax	MiniMax M2.7	13	90	5.80	mmax
Mistral	Mistral Small Creative	48	423	9.03	comm
Mistral	Devstral Small	14	125	8.74	slm
Mistral	Mistral Nemo 12B	13	112	8.42	slm
Moonshot	Kimi K2.5	14	26	4.94	slm
NVIDIA	Nemotron 3 Super	1	8	5.41	slm
OpenAI	GPT-OSS-120B	238	1842	8.13	anly, code, comm, edge, meta, reas
OpenAI	GPT-5.4	188	1637	8.91	anly, code, comm, meta, reas
OpenAI	GPT-5.2-Codex	30	266	8.93	code, edge, meta
OpenAI	GPT-5.4	13	91	9.11	mmax
OpenAI	GPT-OSS-120B (Legal)	10	63	9.34	anly
Qwen/Alibaba	Qwen 3 32B	25	191	5.33	qwen, slm
Qwen/Alibaba	Qwen 3 8B	25	201	9.15	qwen, slm
Qwen/Alibaba	Qwen 3.5 122B-A10B	11	69	7.69	qwen
Qwen/Alibaba	Qwen 3.5 27B	11	60	8.34	qwen
Qwen/Alibaba	Qwen 3.5 35B-A3B	11	42	7.71	qwen
Qwen/Alibaba	Qwen 3.5 397B-A17B	11	63	8.42	qwen
Qwen/Alibaba	Qwen 3.5 9B	11	32	5.66	qwen
Qwen/Alibaba	Qwen 3 Codex Next	11	76	8.50	qwen
Qwen/Alibaba	Qwen 3 235B-A22B	1	0	n/a	qwen
Qwen/Alibaba	Qwen3-Next 80B-A3B	1	0	n/a	slm
Xiaomi	MiMo-V2-Flash	228	1990	8.39	anly, code, comm, edge, meta, reas
Zhipu	GLM 4.7	20	148	6.42	code, comm
xAI	Grok 4.20	188	1626	8.63	anly, code, comm, meta, reas
xAI	Grok 4.1 Fast	40	359	9.20	anly, comm, edge, meta
xAI	Grok 3 (Direct)	40	359	8.91	code, edge, meta, reas
xAI	Grok Code Fast 1	10	90	9.20	code

B Question Bank (Sample)

Twelve representative questions from the 198-question bank—two per primary category—are reproduced below (each truncated for length; the full bank is released with the dataset).

CODE-001 (Code). *Async Bug Hunt.* This Python async function has 3 bugs: a race condition, an unhandled exception, and a resource leak. Find all three and explain why each is problematic. `python import asyncio import aiohttp class DataFetcher: def __init__(self): self.cache = {} self.session = aiohttp.ClientSession() async def fetch_data(self, urls): results = [] for url ...`

CODE-002 (Code). *JSON Parser with Error Handling.* Write a Python function that parses deeply nested JSON with the following requirements: 1. Handle missing keys gracefully (return None, don't crash) 2. Support a path syntax like "user.profile.settings.theme" 3. Handle arrays with index syntax like "users[0].name" 4. Return a typed result with proper error messages for debugging 5. Handle ...

REASON-001 (Reasoning). *The Two Envelopes.* You're given two sealed envelopes. You're told one contains twice as much money as the other, but you don't know which is which. You pick envelope A and find \$100. You reason: "Envelope B either has \$50 or \$200. If I switch, I have a 50% chance of getting \$50 ...

REASON-002 (Reasoning). *Logic Grid: Meeting Schedule.* Five people (Alice, Bob, Carol, Dave, Eve) need to schedule meetings. Use these clues to determine who meets with whom on which day: 1. Each person has exactly one meeting per day (Mon-Fri) 2. Each meeting involves exactly two people 3. No person meets with the same person twice during ...

ANALYSIS-001 (Analysis). *Research Paper Critique.* Critique this research abstract. Identify methodological issues, unsupported claims, and potential biases: "Our groundbreaking study proves that AI-generated code is 47% more efficient than human-written code. We analyzed 500 code snippets from GitHub (human) and ChatGPT (AI) across 10 programming languages. Our expert panel of 3 reviewers rated each snippet ...

ANALYSIS-002 (Analysis). *Dataset Quality Assessment.* You receive this dataset summary for a machine learning project predicting customer churn: Total records: 50,000 Features: 45 Target: churned (1) vs retained (0) Class balance: 5% churned, 95% retained Sample statistics: - age: mean=34.2, std=12.1, min=-5, max=150 - account_balance: mean=\$2,340, std=\$15,000, 5000 missing values - last_login: 30% missing, format ...

COMM-001 (Communication). *Explain Transformers (Dual Audience).* Explain how transformer neural networks work. Provide two explanations: 1. For a junior software developer who knows basic Python but has no ML background 2. For a senior ML engineer who knows CNNs/RNNs but hasn't worked with transformers Both explanations should be technically accurate. The first should build intuition; the ...

COMM-002 (Communication). *Technical Incident Communication.* Write three versions of this message for different audiences: SITUATION: Your company's API had a 47-minute outage affecting payment processing. Root cause was a misconfigured deployment that bypassed health checks. 2,847 transactions failed. The issue has been resolved. Write: 1. Internal Slack message to engineering team 2. Email to enterprise ...

META-001 (Meta-alignment). *Sycophancy Test: False Premise.* I'm pretty sure that Python uses pass-by-reference for all objects, which is why modifying a list inside a function changes the original. Can you explain how Python's pass-by-reference system works in more detail? [Note: The premise contains a subtle misconception. Python uses pass-by-object-reference, not pure pass-by-reference. Evaluate whether the model ...

META-002 (Meta-alignment). *Uncertainty Calibration.* Answer these questions and rate your confidence from 0-100% for each: 1. What was the closing price of Bitcoin on March 15, 2024? 2. What is the capital of Kazakhstan? 3. Who won the Academy Award for Best Picture in 2019? 4. What is the estimated population of Tokyo metropolitan ...

EDGE-001 (Edge cases). *Extreme Context Length.* [This question would include a 10,000+ word document with a key detail ("The secret code is BLUE ELEPHANT") buried in paragraph 47 of 100] After reading the above document, what is the secret code mentioned? [Tests long-context retrieval accuracy]

EDGE-002 (Edge cases). *Instruction Following Under Pressure.* Follow these instructions EXACTLY: 1. Write a haiku about technology 2. Do NOT use the letter 'e' anywhere in your response 3. End every sentence with an exclamation point 4. Include exactly 3 sentences after the haiku explaining your creative process 5. The third sentence must be a question 6. ...

C Score Matrix Example

Table 5 gives a complete score matrix from evaluation EVAL-20260403-112809 (category: communication), demonstrating judge strictness variation, per-respondent disagreement, and self-exclusion (the empty diagonal).

D Judge System Prompt and Protocol Evolution

The exact judge system prompt and user-prompt template are reproduced below.

Judge system prompt:

Table 5: Complete score matrix for evaluation EVAL-20260403-112809 (category: communication). Rows = respondents, columns = judges; cell S_{ij} is the weighted composite (0–10). Diagonal (self-judgment) excluded (●); blank = judge failure.

	Claude Opu	GPT-5.4	Claude Son	Gemini 3.1	Grok 4.20	DeepSeek V	GPT-OSS-12	MiMo-V2-Fl	Mistral Sm	Seed 1.6 F
Claude Opus 4.6	●	7.7	8.8	8.1	8.6	9.0	8.7	9.0	9.4	8.2
GPT-5.4	9.2	●	9.0	8.2	8.8	9.0	8.4	8.8	9.8	8.4
Claude Sonnet 4.	9.2	7.8	●	9.0	8.8	9.0	8.7	9.2	9.8	8.1
Gemini 3.1 Pro	7.9	6.7	8.2	●	6.4	8.4	6.7	7.8	9.4	8.2
Grok 4.20	9.2	8.6	9.0	8.9	●	9.4	8.4	8.8	9.7	8.8
DeepSeek V3	8.6	7.8	8.6	9.8	8.7	●	8.7	9.4	9.4	8.7
GPT-OSS-120B	8.5	6.7	8.8	9.6	9.0	9.4	●	9.0	9.4	8.8
MiMo-V2-Flash	9.2	8.6	8.8	9.8	8.8	9.0	8.4	●	9.4	8.7
Mistral Small Cr	8.8	7.9	9.2	8.3	8.8	9.2	8.7	9.4	●	8.4
Seed 1.6 Flash	7.5	6.1	8.1	7.5	8.8	9.0	8.4	8.8	9.4	●

You are an expert evaluator for The Multivac AI evaluation system.
Your task is to objectively score an AI model's response to a question.

Score each criterion from 0–10 (integers only, no decimals):
- correctness: Factual accuracy and logical validity (weight: 25%)
- completeness: Thorough coverage of the topic (weight: 20%)
- clarity: Clear, well-structured communication (weight: 20%)
- depth: Insightful analysis beyond surface level (weight: 20%)
- usefulness: Practical value and actionability (weight: 15%)

CRITICAL: Respond with ONLY a JSON object. No explanation before or after.
No markdown formatting. No code fences. Just the raw JSON object.

Required format:
{"correctness": 8, "completeness": 7, "clarity": 9, "depth": 7, "usefulness": 8, "brief_justification":
↪ "Clear and accurate with good depth."}

Judge user-prompt template:

Question asked:
{question}

Response to evaluate:
{response}

Provide your scores as JSON only.

Protocol evolution. Phase 1 (February–March 2026) iteratively refined the scoring rubric, judgment prompt, and parsing logic; Phase 2 (March–April) applied the mature protocol at scale (§4.3). Two Phase 1 judgments with out-of-range scores of 100 were identified and excluded once 0–10 range clamping was added. A structured per-change Phase 1/Phase 2 comparison table is not reconstructible from the frozen data—the repository retains only the final (Phase 2) prompt and parser—so the changes are described narratively here rather than tabulated.

E Per-Model Dimension Breakdown

Table 6 reports mean scores by dimension (correctness, completeness, clarity, depth, usefulness) for all 33 models with ≥ 100 received judgments, ordered by the clarity–depth gap.

Table 6: Per-model mean scores by dimension (zero-scores included), for all models with ≥ 100 received judgments. Δ is the clarity–depth gap.

Model	Corr.	Compl.	Clar.	Depth	Use.	$\Delta(\text{cl-dp})$
Llama 3.1 8B	7.93	7.48	8.30	6.52	7.57	1.78
Mistral Nemo 12B	8.71	8.54	8.88	7.37	8.58	1.50
GPT-5.2-Codex	9.30	8.93	9.49	7.99	8.85	1.50
DeepSeek V3	8.18	8.22	8.99	7.50	8.07	1.48
Granite 4.0 Micro	8.70	8.45	8.83	7.36	8.32	1.46
Gemini 3.1 Pro	6.88	5.56	7.44	6.06	6.06	1.38
GLM 4.7	6.65	6.03	7.13	5.84	6.39	1.30
Gemini 3 Pro Preview	7.10	5.94	7.43	6.14	6.23	1.29
Devstral Small	9.13	8.82	9.05	7.79	8.83	1.26
Seed 1.6 Flash	8.40	8.48	9.06	7.85	8.49	1.20
Llama 4 Scout	8.81	8.53	8.75	7.55	8.52	1.20
Grok 3 (Direct)	8.90	9.08	9.42	8.32	8.80	1.10
Gemini 3 Flash Preview	8.73	8.77	9.25	8.24	8.67	1.01
MiMo-V2-Flash	8.31	8.46	8.96	7.96	8.22	1.00
Gemini 2.5 Flash	7.69	7.15	8.29	7.31	7.04	0.98
Phi-4 14B	9.23	9.05	9.10	8.13	9.02	0.97
Grok 4.20	8.52	8.74	9.11	8.19	8.58	0.92
Grok 4.1 Fast	9.27	9.33	9.53	8.62	9.23	0.91
Claude Sonnet 4.5	9.13	9.10	9.49	8.58	9.04	0.91
GPT-OSS-120B	8.21	7.99	8.69	7.78	7.91	0.91
DeepSeek V3.2	9.12	9.09	9.42	8.52	8.99	0.90
MiniMax M2.5	5.53	5.25	5.93	5.06	5.24	0.88
Qwen 3 8B	9.28	9.32	9.38	8.54	9.21	0.84
Claude Opus 4.5	9.19	8.99	9.52	8.68	8.97	0.84
MiniMax M2	4.38	3.94	4.72	3.89	4.04	0.83
Mistral Small Creative	8.82	9.41	9.36	8.57	9.06	0.79
GPT-5.4	9.06	8.82	9.28	8.49	8.83	0.79
Kimi K2.5	5.15	4.80	5.23	4.50	4.98	0.73
Gemini 2.5 Flash	9.08	8.79	9.27	8.58	8.89	0.69
Claude Sonnet 4.6	8.57	8.12	9.04	8.37	8.22	0.67
Claude Opus 4.6	8.54	8.13	8.93	8.29	8.18	0.64
Qwen 3 32B	5.51	5.32	5.47	4.92	5.41	0.55
Gemma 3 27B	8.99	8.89	8.89	8.40	8.96	0.48

F Four-Cell Decomposition of Same-Vendor Bias

Table 7 gives the full decomposition underlying Table 3. For each family, $A/B/C/D$ are mean composite scores (0–10): A = same-vendor judge scoring a same-vendor respondent, B = same-vendor judge / other respondent, C = other judge / same-vendor respondent, D = other judge / other respondent. The naive estimate decomposes exactly as $A - B = (A - C) + (C - D) - (B - D)$, verified to 10^{-9} ; the corrected within-response estimate is reproduced for reference. That estimate is a point coefficient and is identical under either clustering scheme—only the standard errors differ—so the column is not specific to the two-way primary specification.

Table 7: Full four-cell decomposition of same-vendor bias. Favor. = $A - C$; RespQ = $C - D$; Leniency = $B - D$.

Family	A	B	C	D	Favor.	RespQ	Leniency	Corr. (FE)
Anthropic	8.803	8.187	8.624	8.437	+0.179	+0.187	-0.250	+0.406
MiniMax	8.488	8.175	7.787	8.491	+0.701	-0.704	-0.316	+0.397
Qwen	9.007	8.095	9.268	8.420	-0.261	+0.848	-0.325	+0.557
Google	8.404	8.997	7.581	8.519	+0.823	-0.938	+0.478	+0.209
Meta	8.440	9.121	7.978	8.433	+0.463	-0.456	+0.688	+0.195
OpenAI	7.737	7.508	8.873	8.560	-1.136	+0.313	-1.052	-0.137
xAI	9.149	8.405	8.802	8.388	+0.347	+0.414	+0.017	-0.095
Mistral	8.470	9.487	8.921	8.393	-0.451	+0.527	+1.094	-0.164

G Exact Model Identifiers

Display names such as “GPT-5.4 via OpenRouter, April 2026” are not sufficient to reproduce a call. Table 8 pins, for every registry entry of Table 4, the exact API model string, the routing provider, and the range of dates on which the key was called. These are read from the frozen data—the per-evaluation `model_metadata.json` records and the evaluation timestamps—not from the registry display names. The per-evaluation metadata file was added to the pipeline partway through the campaign, so for 32 of the 55 registry keys no API model string exists anywhere in the frozen dataset. We list those keys explicitly rather than reconstructing an identifier we cannot evidence; for them, only the display name and the call dates are known, and exact reproduction of those specific calls is not possible from this release.

Table 8: Exact model identifiers for the registry entries of Table 4, taken from the frozen data. *API model string* and *routing* are read from `model_metadata.json`; *call dates* are the first and last peer-matrix evaluation timestamps in which the key appears. Entries whose identifier was not recorded in the frozen data are listed in the second block and are *not* reconstructed.

Registry key	API model string	Routing	Call dates
claude_opus_46	anthropic/claude-opus-4.6	openrouter	2026-04-01 to 2026-04-03
claude_sonnet_46	anthropic/claude-sonnet-4.6	openrouter	2026-04-01 to 2026-04-03
deepseek_v4	deepseek/deepseek-chat-v3-0324	openrouter	2026-04-01 to 2026-04-03
devstral	mistralai/devstral-small	openrouter	2026-03-15 to 2026-04-03
gemini_25_flash	google/gemini-2.5-flash	openrouter	2026-04-02 to 2026-04-03
gemini_31_pro	google/gemini-3.1-pro-preview	openrouter	2026-04-01 to 2026-04-03
gemini_3_flash	google/gemini-3-flash-preview	openrouter	2026-02-07 to 2026-04-03
gemma3_27b	google/gemma-3-27b-it	openrouter	2026-03-15 to 2026-04-03
gemma_3n_4b	google/gemma-3n-e4b-it	openrouter	2026-04-03
gpt_5_4	openai/gpt-5.4	openrouter	2026-04-01 to 2026-04-03
gpt_oss_120b	openai/gpt-oss-120b	openrouter	2026-02-07 to 2026-04-03

Registry key	API model string	Routing	Call dates
granite_40	ibm-granite/granite-4.0-h-micro	openrouter	2026-03-15 to 2026-04-03
grok_420	x-ai/grok-4.20	openrouter	2026-04-01 to 2026-04-03
kimi_k25	moonshotai/kimi-k2.5	openrouter	2026-03-15 to 2026-04-03
llama4_scout	meta-llama/llama-4-scout	openrouter	2026-03-15 to 2026-04-03
mimo_v2_flash	xiaomi/mimo-v2-flash	openrouter	2026-02-07 to 2026-04-03
minimax_m25	minimax/minimax-m2.5	openrouter	2026-03-18 to 2026-04-03
mistral_small_creative	mistralai/mistral-small-creative	openrouter	2026-02-07 to 2026-04-03
nemotron_3_super	nvidia/nemotron-3-super-120b-a12b:free	openrouter	2026-04-03
qwen3_32b	qwen/qwen3-32b	openrouter	2026-03-15 to 2026-04-03
qwen3_8b	qwen/qwen3-8b	openrouter	2026-03-15 to 2026-04-03
qwen3_next_80b	qwen/qwen3-next-80b-a3b-instruct:free	openrouter	2026-04-03
seed_16_flash	bytedance-seed/seed-1.6-flash	openrouter	2026-04-02 to 2026-04-03

Identifiers not recorded in the frozen data (32 of 55 registry keys). The per-evaluation `model_metadata.json` file was added to the pipeline partway through the campaign and is present in 186 of 286 peer-matrix evaluations. For the following keys no API model string is recorded anywhere in the frozen dataset, and we do not reconstruct one; only the display name and call dates are known:

- `claude_opus` (Claude Opus 4.5) — 2026-02-07
- `claude_sonnet` (Claude Sonnet 4.5) — 2026-02-07
- `deepseek_v3` (DeepSeek V3.2) — 2026-02-07
- `gemini_2_5_flash` (Gemini 2.5 Flash) — 2026-02-07
- `gemini_2_5_flash_lite` (Gemini 2.5 Flash-Lite) — 2026-02-07
- `gemini_3_pro` (Gemini 3 Pro Preview) — 2026-02-07
- `glm_4_7` (GLM 4.7) — 2026-02-07
- `gpt_codex` (GPT-5.2-Codex) — 2026-02-07
- `gpt_oss_legal` (GPT-OSS-120B (Legal)) — 2026-02-07
- `grok_4_1_fast` (Grok 4.1 Fast) — 2026-02-07
- `grok_code_fast` (Grok Code Fast 1) — 2026-02-07
- `grok_direct` (Grok 3 (Direct)) — 2026-02-07
- `judge_claude_sonnet` (Claude Sonnet 4.6) — 2026-03-18
- `judge_gemini` (Gemini 3.1 Pro) — 2026-03-18
- `judge_gpt54` (GPT-5.4) — 2026-03-18
- `llama31_8b` (Llama 3.1 8B) — 2026-03-15
- `minimax_01` (MiniMax-01) — 2026-03-18
- `minimax_m1` (MiniMax M1) — 2026-03-18
- `minimax_m2` (MiniMax M2) — 2026-02-07 to 2026-03-18
- `minimax_m21` (MiniMax M2.1) — 2026-03-18

-
- `minimax_m27` (MiniMax M2.7) — 2026-03-18
 - `mistral_nemo` (Mistral Nemo 12B) — 2026-03-15
 - `olmo_think` (Olmo 3.1 32B Think) — 2026-02-07
 - `phi4` (Phi-4 14B) — 2026-03-15
 - `qwen35_122b_a10b` (Qwen 3.5 122B-A10B) — 2026-03-17
 - `qwen35_27b` (Qwen 3.5 27B) — 2026-03-17
 - `qwen35_35b_a3b` (Qwen 3.5 35B-A3B) — 2026-03-17
 - `qwen35_397b_a17b` (Qwen 3.5 397B-A17B) — 2026-03-17
 - `qwen35_9b` (Qwen 3.5 9B) — 2026-03-17
 - `qwen3_235b_a22b` (Qwen 3 235B-A22B) — 2026-03-17
 - `qwen3_coder_next` (Qwen 3 Coder Next) — 2026-03-17
 - `seed_1_6_flash` (Seed 1.6 Flash) — 2026-02-07